Navigating the R Package Universe

by Spencer Graves, John C. Nash, Julia Silge

Abstract An abstract of less than 150 words.

Introduction

As of our writing, there are more than 12,000 packages on CRAN. R users must approach this abundance of packages with effective strategies to find what they need and choose which packages to invest time in learning how to use. At useR!2017 in Brussels, we organized a contributed session centered on this issue, with three themes in our discussion. The three themes we focused on are **search**, **guidance**, and **unification**. Here, we summarize these important themes, the discussion in our community both at useR!2017 and in the intervening months, and where we can go from here.

Users need options to search R packages, perhaps the content of DESCRIPTION files, documentation files, or other components of R packages. One author (Graves) has worked on the issue of searching for R functions from within R itself in the sos package (Graves et al., 2017), and other options have been built such as RDocumentation.org (Cornelissen, 2018).

Guidance about what package to use for any given task is available from multiple resources for users. R users can turn to long-established resources like CRAN Task Views (reference?), or newer options under current development such as the packagemetrics package (Firke et al., 2018) or the CRANsearcher RStudio addin (Krouse and Calatroni, 2018). One author (Silge) organized a survey before useR about how R users learn about R packages that informed our discussion and is summarized here.

By unification, we largely mean meta-packages or wrappers, packages that call other, related packages for a common set of tasks. With a unified wrapper package, a user only has to learn one API but then can use many different implementations for a certain task. One author (Nash) has been particularly involved in numerical optimization techniques and presented possibilities there and beyond. More generally, as revealed during breakout discussions at useR!2017 and beyond, there are opportunities to merge either packages or their functionality. Such ideas require human cooperation and some give and take in a realm where egos can take precedence over the efficiency of the R ecosystem.

After our main presentation at useR!2017, we broke out into three smaller sessions focused on these three themes. We are encouraged by the engaged attendance and vigorous participation from the community we experienced, and hope to use our community's enthusiasm and ideas to move forward with steps that will improve parts of the R ecosystem.

You can include references in parentheses (R Core Team, 2012), or cite a reference such as R Core Team (2012) in the text.

Search

Guidance

In preparation for this session, one author (Silge) ran a brief online survey in the spring of 2017 to ask R users how they currently discover and learn about R packages. The results of this survey are available in an R package (Silge, 2018) on GitHub. There were 1039 respondents to this survey, which had a single multiple select question on it, "How do you currently discover and learn about R packages?"

Responses to this survey were fielded from R email help lists, local R meetup groups, social media such as Twitter, and affinity groups such as R-Ladies. Figure 1 shows when users responded to the survey. The respondents to this survey overwhelmingly look to social media including blogs and Twitter to learn about R packages, and also make use of general search sites and their personal network.

There were helpful, insightful answers from people contributing to the "other" option. R users use Stack Overflow to learn about R packages, as well as options like CRANberries and crantastic, both of which have RSS feeds that users follow. Other users mentioned learning by reading code on GitHub, and other search websites including rpackages.io.

At useR!2017, after the large contributed session, we broke out into three smaller sessions for discussion and brainstorming. In the breakout session focused on guidance for package choice and package evaluation, we had about 40 participants in our discussion. It was a fruitful discussion and several important themes emerged.

Table 1: Percentage of respondents who chose each answer on survey

How do you currently discover and learn about R packages?	% of respondents
Social media such as blogs, R-bloggers, Twitter, Slack, or GitHub contacts	79.8%
General search websites such as Google and Yahoo	57.0%
Your personal network, such as colleagues and professors	41.6%
Books, textbooks, or journal articles (JSS, JOSS, R-Journal)	31.9%
Conferences, meet-ups, or seminars	24.1%
CRAN Task Views	21.8%
Email lists such as r-help, r-packages, or r-pkg-devel	15.3%
R-specific search websites such as METACRAN or Rdocumentation	11.1%
Other	4.2%
R packages built for search such as the sos package	2.2%
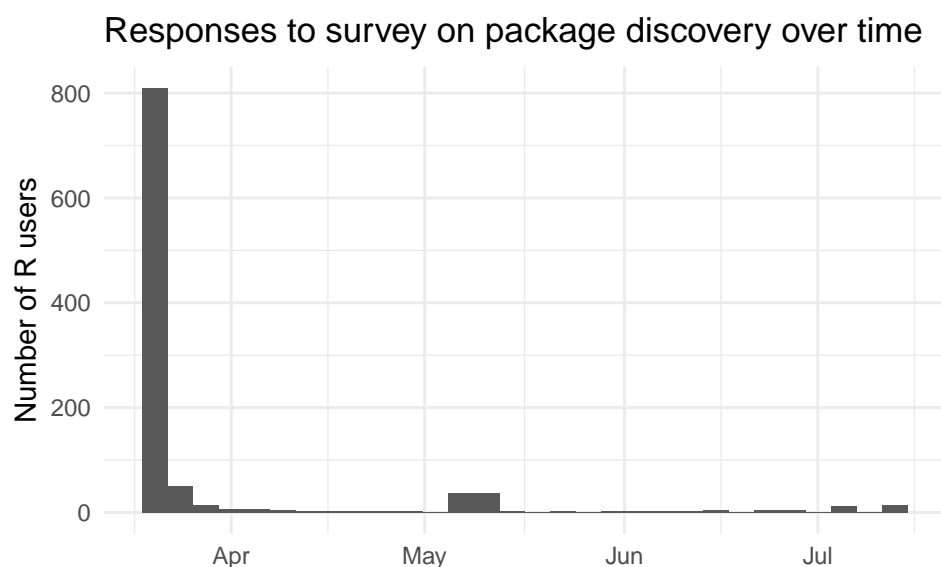
Figure 1: Responses to survey on package discovery during the spring of 2017

Value of personal impact

Participants in this session emphasized how impactful personal relationships can be in how packages are shared and evaluated. Some participants discussed how building local networks of R users may be more important in this effort than top-down, technological solutions. Our survey does show that personal recommendations have been important for many individuals in evaluating R packages. This is yet another area where local user groups can continue to have important impact. Some ways to share this experience more broadly would be online video series or live data analysis, such as those by Sean Taylor and Roger Peng.

CRAN Task Views

Some participants wondered whether the idea of a CRAN Task View is outdated in the current climate with so many packages, and whether it is even possible for one person to maintain one effectively. Others responded that CTVs are all about curation, which is still important, perhaps even more important now. We had at least one CTV maintainer present in our breakout session, and several things were presented as important in order for CTV maintainers to do their jobs:

- Package maintainers should update their NEWS files.
- Package maintainers need to write good documentation.

These are helpful for *all* R users, of course, but also for maintainers of CRAN Task Views. The pkgdown (Wickham and Hesselberth, 2018) package was mentioned as a great way to make documentation visible.

CRAN and *you*

Participants had several ideas about how things are done on CRAN now and adjustments that might be made in the interest of discovering and evaluating packages. One idea that came up several times was the possibility of keywords or tagging for packages. Since useR!2017, the authors have learned that there is support for some tagging architecture for packages on CRAN in the DESCRIPTION file using ACM, JEL, or MSC classifications. For an example of this in action, check out the lfe (Gaure, 2018) package. These are fairly unwieldy lists currently and something like an RStudio addin could be used to navigate them, if they were widely used.

Another desire participants voiced was for more information directly on CRAN, such as the number of downloads for packages. Participants also suggested that vignettes for context-specific tasks like the Bioconductor Workflows would be helpful for package discovery and evaluation, either associated with CRAN or perhaps the *R Journal*. Finally, there was some discussion about whether the very minimal gate-keeping on CRAN was good or bad for the community, although the general feeling was that efforts to keep packages off CRAN would not be positive.

More data, more problems

Some of the package developers at the session wondered why, when R is a data-centric language, developers have such primitive analytics about their users. Issues of user privacy are central here, but there might be opt-in options that could help both package developers and users make better decisions. The idea of a recommender system for R packages was brought up multiple times, perhaps a Tinder for R packages like papr, the Tinder for academic preprints. Both the users and developers present thought that data on package use (instead of package downloads alone) would be helpful in evaluating how important or helpful R packages are. Participants also discussed the possibility of a linter for analysis scripts, similar in concept to linters for code (such as Hester (2017)), that would suggest packages and good practice. Such a linter would necessarily be opinionated, but almost all of the efforts to suggest and evaluate R packages are at some level.

Unification

Summary

Our work on these topics leads us to call for increased respect and value for the work done by local meetup group organizers and individuals who contribute to spreading R knowledge, both online and in their communities. Our survey and discussions show how impactful these community networks are; investing in community building is not something we need do only because of idealism, but because it is effective.

We can also see the importance of continued commitment to growing the skills of package developers across the R ecosystem. Adopting best practices, including writing good documentation, makes this entire challenge better for everyone, from the CRAN Task View maintainer to the new R user.

JOHN AND SPENCER: For full details of *The R Journal* style and information on how to prepare your article for submission, see the Instructions for Authors.

Bibliography

J. Cornelissen. *RDocumentation: Integrate R with 'RDocumentation'*, 2018. URL https://CRAN.R-project.org/package=RDocumentation. R package version 0.8.2. [p1]

S. Firke, B. Krouse, E. Grand, L. Shepherd, W. Ampeh, and H. Frick. *packagemetrics: Collect Metrics on Packages from CRAN, GitHub, and StackOverflow*, 2018. URL https://github.com/ropenscilabs/packagemetrics. R package version 0.0.1.9001. [p1]

S. Gaure. *lfe: Linear Group Fixed Effects*, 2018. URL https://CRAN.R-project.org/package=lfe. [p3]

S. Graves, S. Dorai-Raj, and R. Francois. *sos: Search Contributed R Packages, Sort by Package*, 2017. URL https://CRAN.R-project.org/package=sos. R package version 2.0-0. [p1]

J. Hester. *lintr: A 'Linter' for R Code*, 2017. URL https://CRAN.R-project.org/package=lintr. R package version 1.0.2. [p3]

B. Krouse and A. Calatroni. *CRANsearcher: RStudio Addin for Searching Packages in CRAN Database Based on Keywords*, 2018. URL https://github.com/RhoInc/CRANsearcher. R package version 1.0.0. [p1]

R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2012. URL http://www.R-project.org/. ISBN 3-900051-07-0. [p1]

J. Silge. *packagesurvey: Survey of R Users on Package Discovery*, 2018. URL https://github.com/juliasilge/packagesurvey. R package version 0.1.0. [p1]

H. Wickham and J. Hesselberth. *pkgdown: Make Static HTML Documentation for a Package*, 2018. URL https://CRAN.R-project.org/package=pkgdown. R package version 1.1.0. [p3]

Spencer Graves
EffectiveDefense.org
line 1
line 2
spencer.graves@effectivedefense.org

John C. Nash
University of Ottawa
line 1
line 2
nashjc@uottawa.ca

Julia Silge
Stack Overflow
110 William Street, Floor 28
New York City, NY 10038
julia.silge@gmail.com